

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import xgboost
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
import warnings
warnings.filterwarnings("ignore")
```

```
dataset = pd.read_csv(r"C:\Users\afroz\OneDrive\Desktop\smartbridge\traffic volume.csv")
```

```
dataset.head()
```

```
↗
```

| | holiday | temp | rain | snow | weather | date | Time | traffic_volume |
|---|---------|--------|------|------|---------|------------|----------|----------------|
| 0 | NaN | 288.28 | 0.0 | 0.0 | Clouds | 02-10-2012 | 09:00:00 | 5545 |
| 1 | NaN | 289.36 | 0.0 | 0.0 | Clouds | 02-10-2012 | 10:00:00 | 4516 |
| 2 | NaN | 289.58 | 0.0 | 0.0 | Clouds | 02-10-2012 | 11:00:00 | 4767 |
| 3 | NaN | 290.13 | 0.0 | 0.0 | Clouds | 02-10-2012 | 12:00:00 | 5026 |
| 4 | NaN | 291.14 | 0.0 | 0.0 | Clouds | 02-10-2012 | 13:00:00 | 4918 |

```
dataset.iloc[35764]
```

```
↗
```

| | |
|----------------|------------|
| holiday | NaN |
| temp | 292.95 |
| rain | 0.0 |
| snow | 0.0 |
| weather | Clear |
| date | 28-07-2017 |
| Time | 23:00:00 |
| traffic_volume | 2488 |

Name: 35764, dtype: object

```
dataset.shape
```

```
↗ (48204, 8)
```

```
dataset['holiday'].value_counts()
```

```
↗
```

| | |
|---------------------------|---|
| holiday | |
| Labor Day | 7 |
| Christmas Day | 6 |
| Thanksgiving Day | 6 |
| Martin Luther King Jr Day | 6 |
| New Years Day | 6 |
| Veterans Day | 5 |
| Columbus Day | 5 |
| Memorial Day | 5 |
| Washingtons Birthday | 5 |
| State Fair | 5 |
| Independence Day | 5 |

Name: count, dtype: int64

```
dataset['holiday'] = dataset['holiday'].fillna('None')
```

```
dataset.info()
```

```
↗
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48204 entries, 0 to 48203
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0    holiday     48204 non-null  object
1    temp        48151 non-null  float64
2    rain        48202 non-null  float64
3    snow        48192 non-null  float64
```

```

4  weather      48155 non-null  object
5  date         48204 non-null  object
6  Time         48204 non-null  object
7  traffic_volume 48204 non-null  int64
dtypes: float64(3), int64(1), object(4)
memory usage: 2.9+ MB

```

▼ Handling Missing values

```
dataset.isna().sum()
```

```

↗ holiday      0
  temp        53
  rain         2
  snow        12
  weather      49
  date         0
  Time         0
  traffic_volume 0
dtype: int64

```

```
#Dropping Holiday columns
```

```
# dataset.drop(columns='holiday',inplace=True)
```

```
dataset.head()
```

```

↗
   holiday  temp  rain  snow  weather  date  Time  traffic_volume
0     None  288.28   0.0   0.0   Clouds  02-10-2012  09:00:00         5545
1     None  289.36   0.0   0.0   Clouds  02-10-2012  10:00:00         4516
2     None  289.58   0.0   0.0   Clouds  02-10-2012  11:00:00         4767
3     None  290.13   0.0   0.0   Clouds  02-10-2012  12:00:00         5026
4     None  291.14   0.0   0.0   Clouds  02-10-2012  13:00:00         4918

```

```
dataset['weather'].mode()
```

```

↗ 0    Clouds
   Name: weather, dtype: object

```

```
dataset['weather'].value_counts()
```

```

↗ weather
Clouds      15144
Clear       13383
Mist         5942
Rain        5665
Snow        2875
Drizzle     1818
Haze        1359
Thunderstorm 1033
Fog          912
Smoke        20
Squall        4
Name: count, dtype: int64

```

```
dataset['weather'].fillna('Clouds',inplace=True)
```

```
dataset.sample()
```

```

↗
   holiday  temp  rain  snow  weather  date  Time  traffic_volume
22856     None  276.23  0.42   0.0   Drizzle  28-04-2016  07:00:00         6154

```

```
num_col=['temp','rain','snow']
```

```

for i in num_col:
    dataset[i]=dataset[i].fillna(dataset[i].mean())

```

```
dataset.isna().sum()
```

```
holiday      0
temp         0
rain         0
snow         0
weather      0
date         0
Time         0
traffic_volume 0
dtype: int64
```

```
dataset.describe()
```

```
temp      rain      snow  traffic_volume
count  48204.000000  48204.000000  48204.000000  48204.000000
mean    281.205351    0.334278    0.000222    3259.818355
std     13.336338    44.789133    0.008168    1986.860670
min      0.000000    0.000000    0.000000     0.000000
25%     272.180000    0.000000    0.000000    1193.000000
50%     282.429000    0.000000    0.000000    3380.000000
75%     291.800000    0.000000    0.000000    4933.000000
max     310.070000    9831.300000    0.510000    7280.000000
```

✓ Label Encoding weather and holiday Columns

This code applies Label Encoding to the weather and holiday columns in the dataset. Each unique category in these columns is converted into an integer value.

weather: Converts weather types like Clear, Clouds, etc., into numbers.

holiday: Converts holiday labels (like None, New Year's Day, etc.) into numbers.

Helps prepare categorical features for machine learning models.

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
dataset['weather'] = le.fit_transform(dataset['weather'])
dataset['holiday'] = le.fit_transform(dataset['holiday'])
```

```
dataset.head(10)
```

```
holiday  temp  rain  snow  weather  date      Time  traffic_volume
0       7  288.28   0.0   0.0       1  02-10-2012  09:00:00      5545
1       7  289.36   0.0   0.0       1  02-10-2012  10:00:00      4516
2       7  289.58   0.0   0.0       1  02-10-2012  11:00:00      4767
3       7  290.13   0.0   0.0       1  02-10-2012  12:00:00      5026
4       7  291.14   0.0   0.0       1  02-10-2012  13:00:00      4918
5       7  291.72   0.0   0.0       0  02-10-2012  14:00:00      5181
6       7  293.17   0.0   0.0       0  02-10-2012  15:00:00      5584
7       7  293.86   0.0   0.0       0  02-10-2012  16:00:00      6015
8       7  294.14   0.0   0.0       1  02-10-2012  17:00:00      5791
9       7  293.10   0.0   0.0       1  02-10-2012  18:00:00      4770
```

✓ Data Visualization

```
numeric_data = dataset.select_dtypes(include='number')
numeric_data.corr()
```



| | holiday | temp | rain | snow | weather | traffic_volume |
|----------------|-----------|-----------|-----------|-----------|-----------|----------------|
| holiday | 1.000000 | -0.000472 | 0.000066 | 0.000432 | -0.004328 | 0.018676 |
| temp | -0.000472 | 1.000000 | 0.009070 | -0.019758 | -0.033559 | 0.130034 |
| rain | 0.000066 | 0.009070 | 1.000000 | -0.000090 | 0.009542 | 0.004714 |
| snow | 0.000432 | -0.019758 | -0.000090 | 1.000000 | 0.036662 | 0.000735 |
| weather | -0.004328 | -0.033559 | 0.009542 | 0.036662 | 1.000000 | -0.040035 |
| traffic_volume | 0.018676 | 0.130034 | 0.004714 | 0.000735 | -0.040035 | 1.000000 |

correltion

```
# Set the plot size and style
plt.figure(figsize=(10, 10))
sns.set(style="white")
numeric_data = dataset.select_dtypes(include='number')
mask = np.triu(np.ones_like(numeric_data.corr(), dtype=bool))
# Create a custom diverging colormap
cmap = sns.diverging_palette(220, 20, as_cmap=True)


# Draw the heatmap
sns.heatmap(numeric_data.corr(),
            mask=mask,
            cmap=cmap,
            annot=True,      # Show correlation values
            fmt=".2f",       # 2 decimal places
            square=True,
            linewidths=.5,
            cbar_kws={"shrink": 0.8},
            annot_kws={"size": 10, "color": "black"})

# Title
plt.title("Feature Correlation Heatmap", fontsize=16, fontweight='bold')

plt.tight_layout()
plt.show()
```



Feature Correlation Heatmap



Pair Plot:

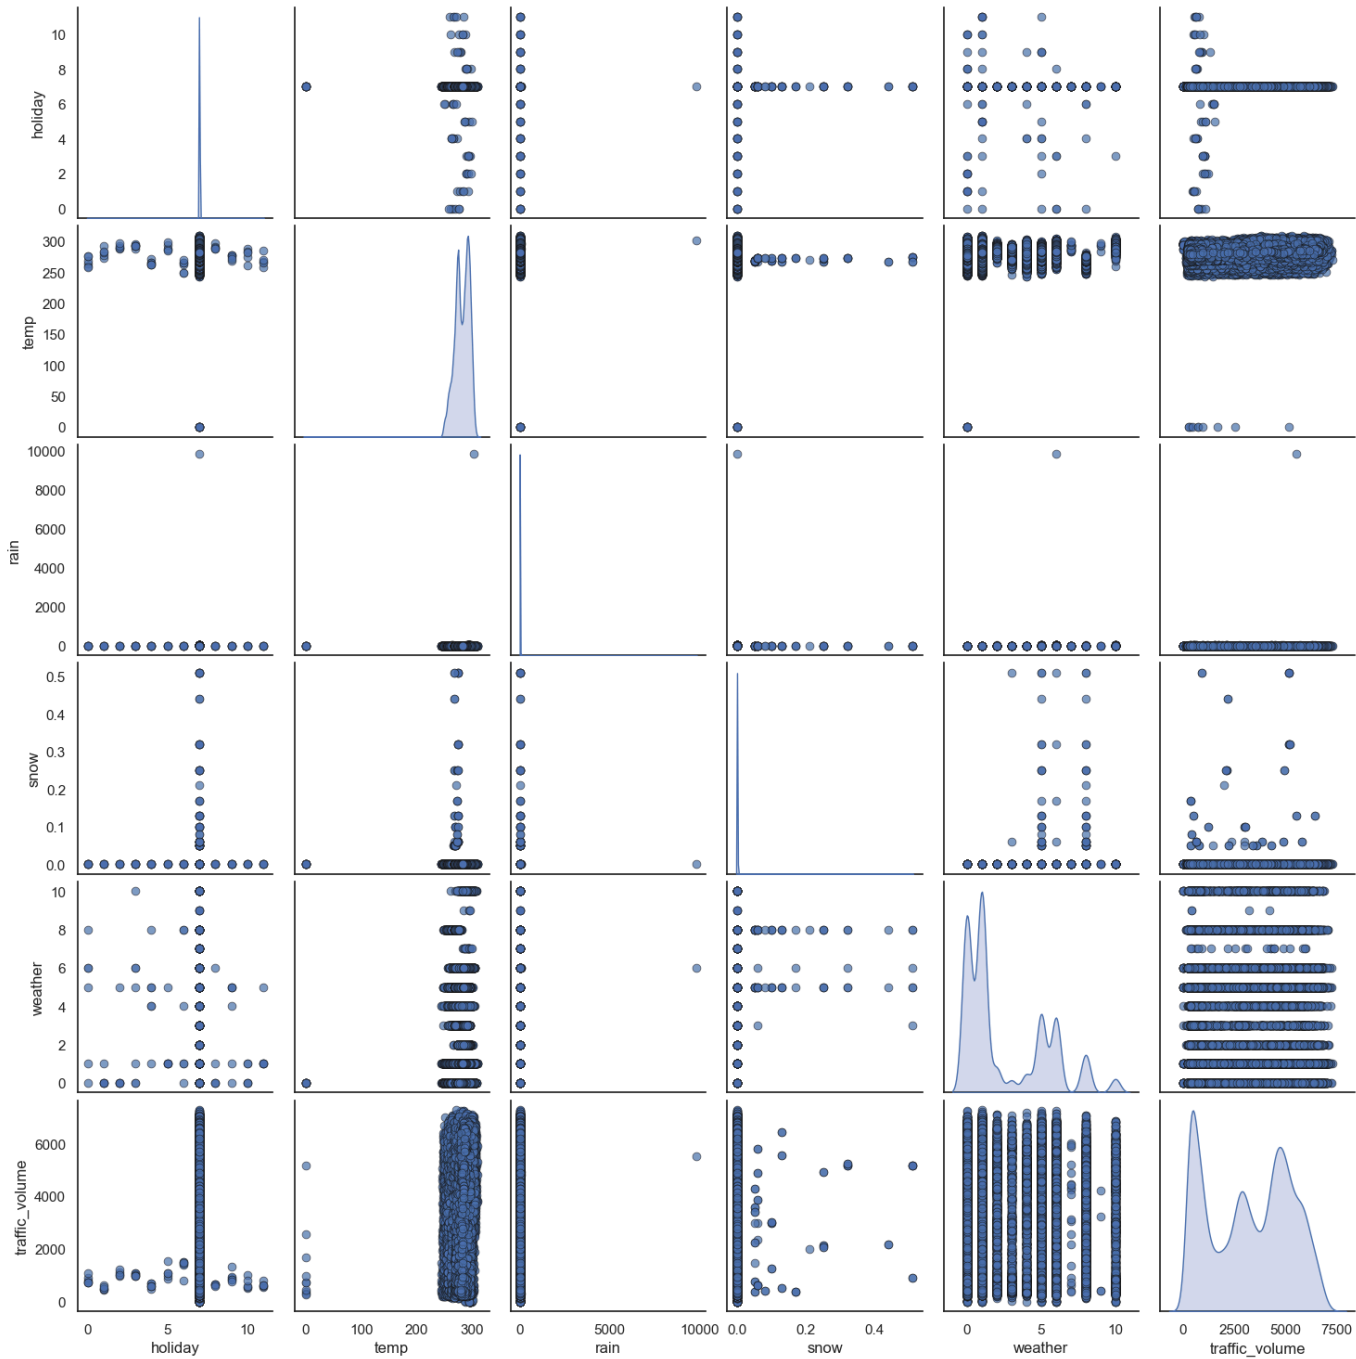
```
# Create pairplot
sns.pairplot(dataset,
              diag_kind='kde',          # Smooth KDE on diagonal

              plot_kws={'alpha': 0.7, 's': 40, 'edgecolor': 'k'}, # Scatter settings
              height=2.5)              # Control size of plots

plt.suptitle("Pairplot of Features", fontsize=16, fontweight='bold', y=1.02)
plt.show()
```



Pairplot of Features

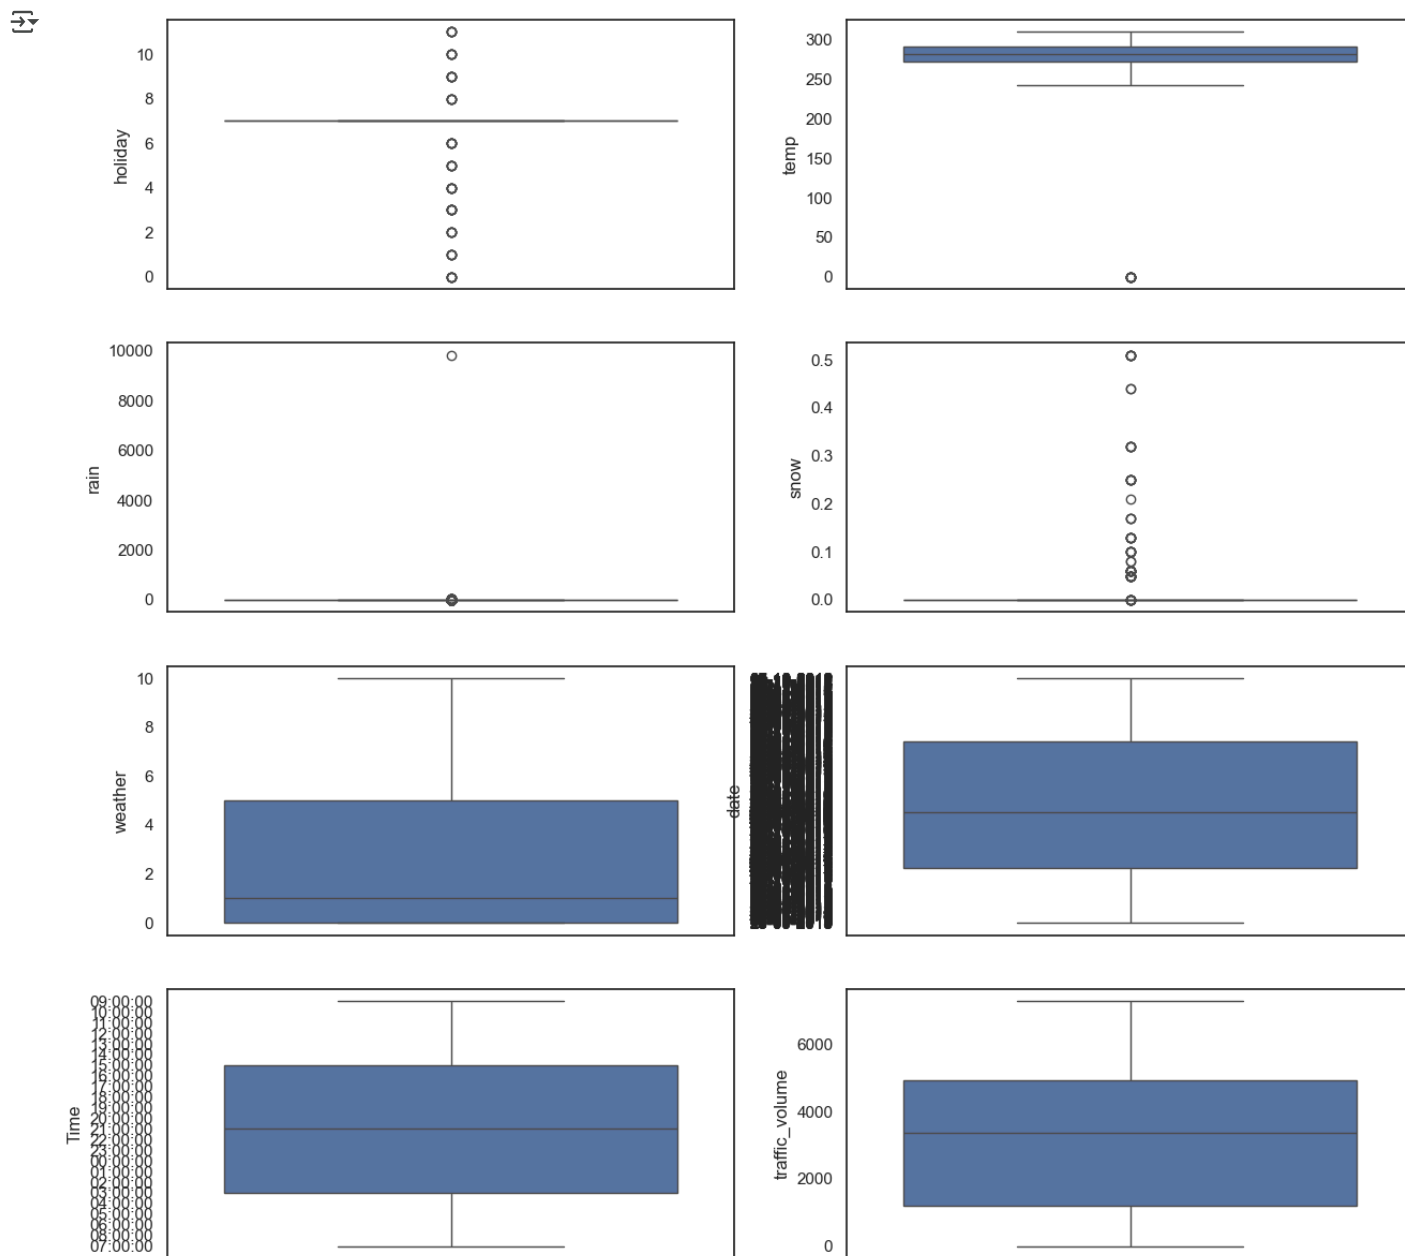


Boxplot

```

fig, axes = plt.subplots(4, 2, figsize=(15, 15))
axes = axes.flatten()
i=0
for c in dataset:
    sns.boxplot(dataset[c], ax=axes[i])
    i+=1

```



Start coding or [generate](#) with AI.

```
#splitting the date column into year,month,day
dataset[['day','month','year']] = dataset['date'].str.split('-',expand=True)
```

```
#splitting the date column into hours minutes seconds
dataset[['hours','minutes','seconds']] = dataset['Time'].str.split(":",expand=True)
```

```
# Dropping date Time columns
dataset.drop(columns=['date','Time'],inplace=True)
```

```
dataset.head()
```

```
↗
```

| | holiday | temp | rain | snow | weather | traffic_volume | day | month | year | hours | minutes | seconds |
|---|---------|--------|------|------|---------|----------------|-----|-------|------|-------|---------|---------|
| 0 | 7 | 288.28 | 0.0 | 0.0 | 1 | 5545 | 02 | 10 | 2012 | 09 | 00 | 00 |
| 1 | 7 | 289.36 | 0.0 | 0.0 | 1 | 4516 | 02 | 10 | 2012 | 10 | 00 | 00 |
| 2 | 7 | 289.58 | 0.0 | 0.0 | 1 | 4767 | 02 | 10 | 2012 | 11 | 00 | 00 |
| 3 | 7 | 290.13 | 0.0 | 0.0 | 1 | 5026 | 02 | 10 | 2012 | 12 | 00 | 00 |
| 4 | 7 | 291.14 | 0.0 | 0.0 | 1 | 4918 | 02 | 10 | 2012 | 13 | 00 | 00 |

```
#Splitting the Dataset into Dependent and Independent variable
x = dataset.drop('traffic_volume',axis=1)
y = dataset['traffic_volume']
```

```
x.head()
```

```
↗
```

| | holiday | temp | rain | snow | weather | day | month | year | hours | minutes | seconds |
|---|---------|--------|------|------|---------|-----|-------|------|-------|---------|---------|
| 0 | 7 | 288.28 | 0.0 | 0.0 | 1 | 02 | 10 | 2012 | 09 | 00 | 00 |
| 1 | 7 | 289.36 | 0.0 | 0.0 | 1 | 02 | 10 | 2012 | 10 | 00 | 00 |
| 2 | 7 | 289.58 | 0.0 | 0.0 | 1 | 02 | 10 | 2012 | 11 | 00 | 00 |
| 3 | 7 | 290.13 | 0.0 | 0.0 | 1 | 02 | 10 | 2012 | 12 | 00 | 00 |
| 4 | 7 | 291.14 | 0.0 | 0.0 | 1 | 02 | 10 | 2012 | 13 | 00 | 00 |

```
# column_names = ['holiday', 'temp', 'rain', 'snow', 'weather', 'day', 'month', 'year',
#                  'hours', 'minutes', 'seconds']
```

```
x.shape,y.shape
```

```
↗ ((48204, 11), (48204,))
```

```
y.head()
```

```
↗
```

| | |
|---|------|
| 0 | 5545 |
| 1 | 4516 |
| 2 | 4767 |
| 3 | 5026 |
| 4 | 4918 |

Name: traffic_volume, dtype: int64

Start coding or [generate](#) with AI.

Feature Scaling

```
name=x.columns
name
```

```
↗ Index(['holiday', 'temp', 'rain', 'snow', 'weather', 'day', 'month', 'year',
        'hours', 'minutes', 'seconds'],
        dtype='object')
```

```
from sklearn.preprocessing import scale
```



```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
x = scaler.fit_transform(x)
x = pd.DataFrame(x, columns=name)
```

```
x.head()
```

```
↵
```

| | holiday | temp | rain | snow | weather | day | month | year | hours | minutes | seconds |
|---|----------|----------|-----------|-----------|-----------|-----------|---------|-----------|-----------|---------|---------|
| 0 | 0.015856 | 0.530485 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.02758 | -1.855294 | -0.345548 | 0.0 | 0.0 |
| 1 | 0.015856 | 0.611467 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.02758 | -1.855294 | -0.201459 | 0.0 | 0.0 |
| 2 | 0.015856 | 0.627964 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.02758 | -1.855294 | -0.057371 | 0.0 | 0.0 |
| 3 | 0.015856 | 0.669205 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.02758 | -1.855294 | 0.086718 | 0.0 | 0.0 |
| 4 | 0.015856 | 0.744939 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.02758 | -1.855294 | 0.230807 | 0.0 | 0.0 |

✓ Splitting the Data into Train and Test

```
# for i in range(1,100):
#     x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=i)
#     model=RandomForestRegressor()
#     model.fit(x_train,y_train)
#     # print(f'{i}-- {classification_report(y_test, y_pred)}')
#     print(f'{i}-- {r2_score(y_train,train3)}')
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=42)
```

Hyper Tuning

```
# from sklearn.ensemble import RandomForestRegressor
# from sklearn.model_selection import GridSearchCV

# # Define parameter grid
# param_grid = {
#     'n_estimators': [100, 200],
#     'max_depth': [10, 20, None],
#     'min_samples_split': [2, 5],
#     'min_samples_leaf': [1, 2]
# }

# # Initialize the model
# rf = RandomForestRegressor(random_state=42)

# # Grid Search with cross-validation
# grid_search = GridSearchCV(estimator=rf, param_grid=param_grid,
#                             cv=3, scoring='r2', n_jobs=-1, verbose=1)

# grid_search.fit(x_train, y_train)

# # Best model
# best_rf = grid_search.best_estimator_
# print("Best Parameters:", grid_search.best_params_)

# y_pred_rf=best_rf.predict(x_train)
# r2_=r2_score(y_train,y_pred_rf)
# print(r2_)

# y_pred_rf=best_rf.predict(x_test)
# r2_=r2_score(y_test,y_pred_rf)
# print(r2_)
```

✓ Training and Testing and Model

```
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
```

```
lr = LinearRegression()
Dtree = DecisionTreeRegressor()
Rand = RandomForestRegressor()
svr = SVR()
XGB = xgboost.XGBRFRegressor()
```

```
# Fit the models with x_tain and y_train
lr.fit(x_train,y_train)
Dtree.fit(x_train,y_train)
Rand.fit(x_train,y_train)
svr.fit(x_train,y_train)
XGB.fit(x_train,y_train)
```

XGBRFRegressor

XGBRFRegressor(base_score=None, booster=None, callbacks=None, colsample_bylevel=None, colsample_bytree=None, device=None, early_stopping_rounds=None, enable_categorical=False, eval_metric=None, feature_types=None, feature_weights=None, gamma=None, grow_policy=None, importance_type=None, interaction_constraints=None, max_bin=None, max_cat_threshold=None, max_cat_to_onehot=None, max_delta_step=None, max_depth=None, max_leaves=None, min_child_weight=None, missing=nan, monotone_constraints=None, multi_strategy=None, n_estimators=None, n_jobs=None, num_parallel_tree=None, objective='reg:squarederror',

```
x_test.shape
```

```
(9641, 11)
```

```
#predict the y_tain value and caculate the accuracy
train1 = lr.predict(x_train)
train2 = Dtree.predict(x_train)
train3 = Rand.predict(x_train)
train4 = svr.predict(x_train)
train5 = XGB.predict(x_train)
```

```
#x
```

| | holiday | temp | rain | snow | weather | day | month | year | hours | minutes | seconds |
|-------|----------|----------|-----------|-----------|-----------|-----------|----------|-----------|-----------|---------|---------|
| 0 | 0.015856 | 0.530485 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.027580 | -1.855294 | -0.345548 | 0.0 | 0.0 |
| 1 | 0.015856 | 0.611467 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.027580 | -1.855294 | -0.201459 | 0.0 | 0.0 |
| 2 | 0.015856 | 0.627964 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.027580 | -1.855294 | -0.057371 | 0.0 | 0.0 |
| 3 | 0.015856 | 0.669205 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.027580 | -1.855294 | 0.086718 | 0.0 | 0.0 |
| 4 | 0.015856 | 0.744939 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.027580 | -1.855294 | 0.230807 | 0.0 | 0.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 48199 | 0.015856 | 0.168313 | -0.007463 | -0.027235 | -0.566452 | 1.635058 | 0.733478 | 1.313958 | 1.095340 | 0.0 | 0.0 |
| 48200 | 0.015856 | 0.116574 | -0.007463 | -0.027235 | -0.566452 | 1.635058 | 0.733478 | 1.313958 | 1.239428 | 0.0 | 0.0 |
| 48201 | 0.015856 | 0.114324 | -0.007463 | -0.027235 | 2.666935 | 1.635058 | 0.733478 | 1.313958 | 1.383517 | 0.0 | 0.0 |
| 48202 | 0.015856 | 0.066334 | -0.007463 | -0.027235 | -0.566452 | 1.635058 | 0.733478 | 1.313958 | 1.527606 | 0.0 | 0.0 |
| 48203 | 0.015856 | 0.068584 | -0.007463 | -0.027235 | -0.566452 | 1.635058 | 0.733478 | 1.313958 | 1.671695 | 0.0 | 0.0 |

48204 rows × 11 columns

```
#dataset.iloc[35764]
```

```
holiday          7
temp            292.95
```

```

rain            0.0
snow            0.0
weather         0
traffic_volume  2488
day             28
month           07
year            2017
hours           23
minutes         00
seconds         00
Name: 35764, dtype: object

```

```
x_train.head()
```

```

↗

```

| | holiday | temp | rain | snow | weather | day | month | year | hours | minutes | seconds |
|-------|----------|-----------|-----------|-----------|-----------|-----------|-----------|----------|-----------|---------|---------|
| 35764 | 0.015856 | 0.880659 | -0.007463 | -0.027235 | -0.925717 | 1.405775 | 0.145275 | 0.785749 | 1.671695 | 0.0 | 0.0 |
| 31011 | 0.015856 | -0.602747 | -0.007463 | -0.027235 | -0.925717 | -0.199205 | -1.325232 | 0.785749 | -0.201459 | 0.0 | 0.0 |
| 28019 | 0.015856 | 0.070833 | -0.007463 | -0.027235 | -0.925717 | -0.428488 | 1.321682 | 0.257541 | 0.086718 | 0.0 | 0.0 |
| 33195 | 0.015856 | -0.177363 | -0.007463 | -0.027235 | 1.229874 | 1.635058 | -0.737029 | 0.785749 | -0.201459 | 0.0 | 0.0 |
| 22348 | 0.015856 | -1.005935 | -0.007463 | -0.027235 | -0.925717 | -0.772412 | -0.737029 | 0.257541 | -1.498258 | 0.0 | 0.0 |

```
y_train
```

```

↗

```

| | |
|-------|------|
| 35764 | 2488 |
| 31011 | 4395 |
| 28019 | 4513 |
| 33195 | 3489 |
| 22348 | 751 |
| ... | |
| 11284 | 5761 |
| 44732 | 4799 |
| 38158 | 5139 |
| 860 | 2057 |
| 15795 | 4385 |

Name: traffic_volume, Length: 38563, dtype: int64

Model Evaluation for X_train data

```

# Predict and evaluate using R², MAE, MSE
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

```

```

print("Linear regression", r2_score(y_train, train1))
print("Decision Tree Regressor", r2_score(y_train, train2))
print("Random Forest Regressor", r2_score(y_train, train3))
print("SVR", r2_score(y_train, train4))
print("XGB", r2_score(y_train, train5))

```

```

↗

```

Linear regression 0.13225712655997102
Decision Tree Regressor 1.0
Random Forest Regressor 0.9779024664396229
SVR 0.24187500547247331
XGB 0.7786684036254883

```

print("Linear regression", mean_squared_error(y_train, train1))
print("Decision Tree Regressor", mean_squared_error(y_train, train2))
print("Random Forest Regressor", mean_squared_error(y_train, train3))
print("SVR", mean_squared_error(y_train, train4))
print("XGB", mean_squared_error(y_train, train5))

```

```

↗

```

Linear regression 3424093.4137532604
Decision Tree Regressor 0.0
Random Forest Regressor 87196.35901395639
SVR 2991543.7856290694
XGB 873369.25

```

print("Linear regression", mean_absolute_error(y_train, train1))
print("Decision Tree Regressor", mean_absolute_error(y_train, train2))
print("Random Forest Regressor", mean_absolute_error(y_train, train3))
print("SVR", mean_absolute_error(y_train, train4))
print("XGB", mean_absolute_error(y_train, train5))

```

```

Linear regression 1640.2171288638979
Decision Tree Regressor 0.0
Random Forest Regressor 185.25322511215415
SVR 1509.2879969637522
XGB 639.703857421875

```

```

print("Linear regression", np.sqrt(mean_squared_error(y_train, train1)))
print("Decision Tree Regressor", np.sqrt(mean_squared_error(y_train, train2)))
print("Random Forest Regressor", np.sqrt(mean_squared_error(y_train, train3)))
print("SVR", np.sqrt(mean_squared_error(y_train, train4)))
print("XGB", np.sqrt(mean_squared_error(y_train, train5)))

```

```

Linear regression 1850.430602252692
Decision Tree Regressor 0.0
Random Forest Regressor 295.29029617303104
SVR 1729.6079861139256
XGB 934.5422676369432

```

✓ Model Evaluation for X_test data

```

# predict the y_test value and caculate the accuracy
test1 = lr.predict(x_test)
test2 = Dtree.predict(x_test)
test3 = Rand.predict(x_test)
test4 = svr.predict(x_test)
test5 = XGB.predict(x_test)

```

```

print("Linear regression", r2_score(y_test, test1))
print("Decision Tree Regressor", r2_score(y_test, test2))
print("Random Forest Regression", r2_score(y_test, test3))
print("SVR", r2_score(y_test, test4))
print("XGB", r2_score(y_test, test5))

```

```

Linear regression 0.1389494912742577
Decision Tree Regressor 0.7169825008411866
Random Forest Regression 0.8430752236706374
SVR 0.2448235297268001
XGB 0.7832188606262207

```

```

print("Linear regression", mean_squared_error(y_test, test1))
print("Decision Tree Regressor", mean_squared_error(y_test, test2))
print("Random Forest Regressor", mean_squared_error(y_test, test3))
print("SVR", mean_squared_error(y_test, test4))
print("XGB", mean_squared_error(y_test, test5))

```

```

Linear regression 3404174.861301238
Decision Tree Regressor 1118913.520381703
Random Forest Regressor 620404.2309739654
SVR 2985600.414724397
XGB 857047.125

```

```

print("Linear regression", mean_absolute_error(y_test, test1))
print("Decision Tree Regressor", mean_absolute_error(y_test, test2))
print("Random Forest Regressor", mean_absolute_error(y_test, test3))
print("SVR", mean_absolute_error(y_test, test4))
print("XGB", mean_absolute_error(y_test, test5))

```

```

Linear regression 1638.7989252319232
Decision Tree Regressor 555.1361891919926
Random Forest Regressor 496.8264121979048
SVR 1510.833149986356
XGB 632.3110961914062

```

```

print("Linear regression", np.sqrt(mean_squared_error(y_test, test1)))
print("Decision Tree Regressor", np.sqrt(mean_squared_error(y_test, test2)))
print("Random Forest Regressor", np.sqrt(mean_squared_error(y_test, test3)))
print("SVR", np.sqrt(mean_squared_error(y_test, test4)))
print("XGB", np.sqrt(mean_squared_error(y_test, test5)))

```

```

Linear regression 1845.0406123717814
Decision Tree Regressor 1057.7870865073476
Random Forest Regressor 787.6574325009352
SVR 1727.889005325399
XGB 925.7683970626779

```

```
x.head()
```

| | holiday | temp | rain | snow | weather | day | month | year | hours | minutes | seconds |
|---|----------|----------|-----------|-----------|-----------|-----------|---------|-----------|-----------|---------|---------|
| 0 | 0.015856 | 0.530485 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.02758 | -1.855294 | -0.345548 | 0.0 | 0.0 |
| 1 | 0.015856 | 0.611467 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.02758 | -1.855294 | -0.201459 | 0.0 | 0.0 |
| 2 | 0.015856 | 0.627964 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.02758 | -1.855294 | -0.057371 | 0.0 | 0.0 |
| 3 | 0.015856 | 0.669205 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.02758 | -1.855294 | 0.086718 | 0.0 | 0.0 |
| 4 | 0.015856 | 0.744939 | -0.007463 | -0.027235 | -0.566452 | -1.574903 | 1.02758 | -1.855294 | 0.230807 | 0.0 | 0.0 |

```
dataset.head()
```

| | holiday | temp | rain | snow | weather | traffic_volume | day | month | year | hours | minutes | seconds |
|---|---------|--------|------|------|---------|----------------|-----|-------|------|-------|---------|---------|
| 0 | 7 | 288.28 | 0.0 | 0.0 | 1 | 5545 | 02 | 10 | 2012 | 09 | 00 | 00 |
| 1 | 7 | 289.36 | 0.0 | 0.0 | 1 | 4516 | 02 | 10 | 2012 | 10 | 00 | 00 |
| 2 | 7 | 289.58 | 0.0 | 0.0 | 1 | 4767 | 02 | 10 | 2012 | 11 | 00 | 00 |
| 3 | 7 | 290.13 | 0.0 | 0.0 | 1 | 5026 | 02 | 10 | 2012 | 12 | 00 | 00 |
| 4 | 7 | 291.14 | 0.0 | 0.0 | 1 | 4918 | 02 | 10 | 2012 | 13 | 00 | 00 |

```
Rand.predict([[7, 289.28, 0.0, 0.0, 1, 2, 10, 2012, 10, 0, 0]])
```

```
array([1600.02])
```

Random Forest is very Less when compared with other models, so Saving the Random forest model and Deploying.

In Random Forest Best result:

- ✓ No overfitting: A small gap (97% → 84%) means the model isn't memorizing, it's learning.
- ✓ Good generalization: Predicts traffic volume on unseen data quite accurately.
- ✓ Feature processing + model choice worked well.

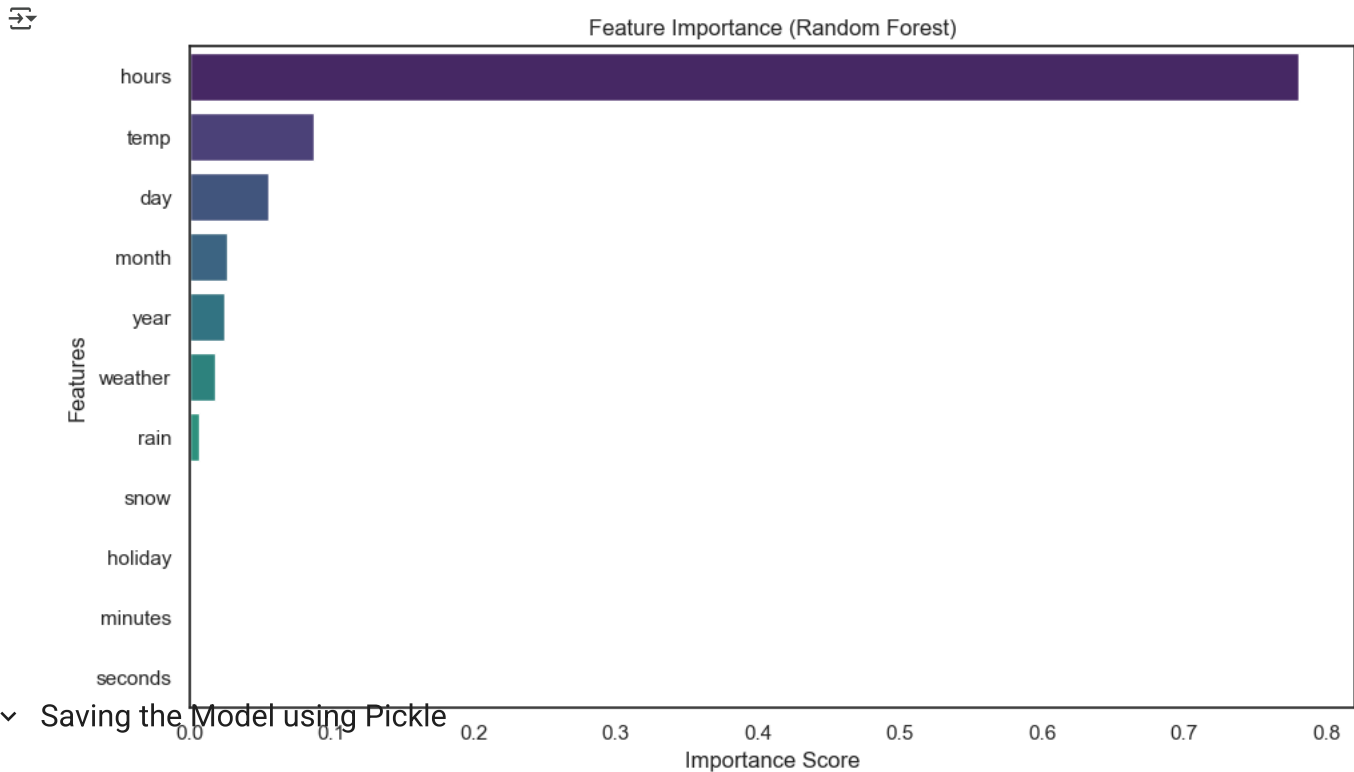
✓ Feature importances

```
Rand.feature_importances_
```

```
array([1.19496933e-05, 8.71820215e-02, 6.91457975e-03, 5.97033853e-05,
       1.77991664e-02, 5.57887871e-02, 2.66752344e-02, 2.47915670e-02,
       7.80776991e-01, 0.00000000e+00, 0.00000000e+00])
```

```
# Get feature importances
importances = Rand.feature_importances_
feature_names = x.columns
feature_df = pd.DataFrame({'Feature': feature_names, 'Importance': importances})
feature_df = feature_df.sort_values(by='Importance', ascending=False)
```

```
# Plot feature importance
plt.figure(figsize=(10, 6))
sns.barplot(data=feature_df, x='Importance', y='Feature', palette='viridis')
plt.title('Feature Importance (Random Forest)')
plt.xlabel('Importance Score')
plt.ylabel('Features')
plt.tight_layout()
plt.show()
```



✓ Saving the Model using Pickle

```
import pickle

# # Assume 'model' is your trained model
# filename = 'finalized_model.pkl'

# # Save the model to disk
# with open(filename, 'wb') as file:
#     pickle.dump(Rand, file)

# with open(filename, 'rb') as file:
#     loaded_model = pickle.load(file)

# # Now you can use it to make predictions
# predictions = loaded_model.predict(x_test)
```